
Mastering System Design

Section 3: Protocols

— Networking & Communication —

Section Agenda

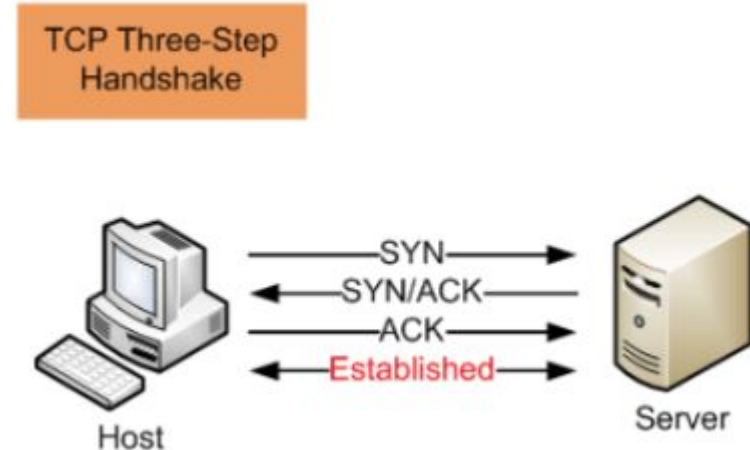
1. Lecture 1: TCP & UDP - Understanding the Basics
2. Lecture 2: HTTP - The Backbone of the Web
3. Lecture 3: REST & RESTfulness - API Design Principles
4. Lecture 4: Real-Time Communication Protocols
5. Lecture 5: Modern API Protocols - Beyond REST
6. Lecture 6: Summary & Practical Applications

TCP & UDP

Protocols


What is Transmission Control Protocol (TCP)?

- Connection-oriented protocol
- Reliable, ordered, and error-checked communication
- Ensures data reaches the destination correctly



What is User Datagram Protocol (UDP)?

- Connectionless protocol
- Faster, but no delivery guarantees
- No retransmission of lost packets





Key Differences Between TCP & UDP

- Reliability: TCP (Yes), UDP (No)
- Speed: TCP (Slower), UDP (Faster)
- Connection Type: TCP (Connection-oriented), UDP (Connectionless)

Feature	TCP (Transmission Control Protocol)	UDP (User Datagram Protocol)
Reliability	✅ Reliable (ensures data delivery)	❌ Unreliable (no guarantee of delivery)
Speed	⌚ Slower (due to error checking & retransmission)	⚡ Faster (no retransmission overhead)
Connection Type	🔗 Connection-oriented (establishes a connection before communication)	📠 Connectionless (sends data without setup)
Ordering	✅ Ensures packets arrive in order	❌ No guarantee of packet order
Error Handling	✅ Built-in error checking & retransmission	❌ Minimal error checking, no retransmission
Overhead	📦 High (due to handshaking, sequencing, and acknowledgments)	📦 Low (minimal protocol overhead)
Use Cases	🌐 Web browsing (HTTP/HTTPS), File transfers (FTP, SFTP), Email (SMTP, IMAP, POP3), Database communication	📺 Video streaming (YouTube, Netflix), 🎮 Online gaming, 📞 VoIP calls (Skype, Zoom), 🌐 DNS lookups

Use Cases: When to Use TCP vs. UDP

- **When to use TCP:** If data integrity is critical, TCP is the way to go.
 - Web browsing (HTTP, HTTPS)
 - File transfers (FTP, SFTP)
 - Email (SMTP, IMAP, POP3)
 - Database communication
- **When to use UDP:** UDP is preferred when speed is more important than reliability
 - Video streaming (YouTube, Netflix)
 - Online gaming (Multiplayer Games)
 - VoIP calls (Skype, Zoom)
 - DNS lookups (Domain Name System)



Interview Questions: TCP & UDP

- Basic Questions:
 - What is the difference between TCP and UDP?
 - Why is TCP considered a reliable protocol?
 - When would you choose UDP over TCP?
- Technical & Deep-Dive:
 - How does TCP ensure reliable data transmission?
 - What is the Three-Way Handshake in TCP?
 - How does UDP handle packet loss?
- Use Case & System Design:
 - Which protocol would you use for a real-time multiplayer game? Why?
 - How does TCP handle congestion control?
 - Why is DNS implemented over UDP instead of TCP?
 - If a video streaming service is buffering too much, would switching from TCP to UDP help?

Pro Tip: Be prepared to compare TCP & UDP in real-world scenarios and discuss trade-offs in speed vs. reliability!

Summary & Key Takeaways


- TCP is reliable but slower; UDP is fast but unreliable.
- TCP is best for web pages, file transfers, and email.
- UDP is best for real-time applications like video calls and gaming.
- System design must balance speed, reliability, and user experience.
- What's Next:
 - HTTP - The Backbone of the Web

HTTP - The Backbone of the Web

— Protocols —

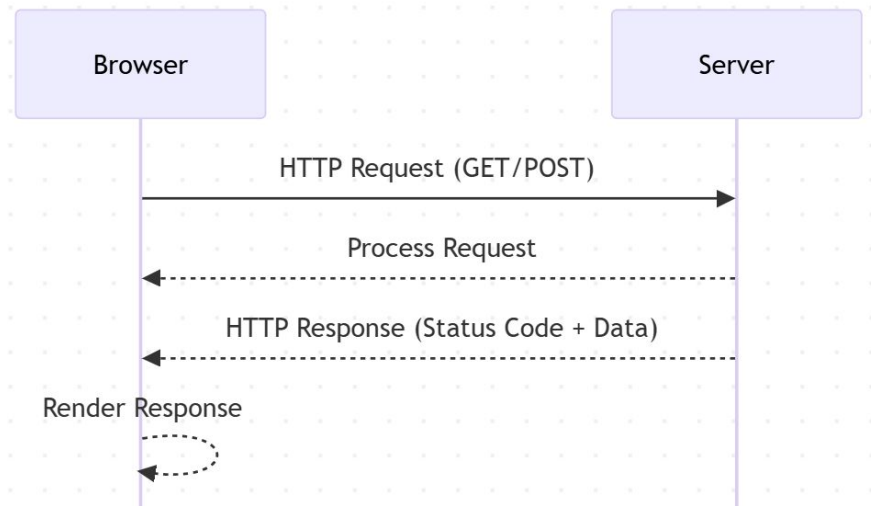
Introduction to HTTP

- What is HTTP?
 - Stands for HyperText Transfer Protocol
 - The foundation of web communication
 - Defines rules for requesting and transferring resources (e.g., web pages, APIs, images)
 - Works over TCP/IP (Port 80 for HTTP, Port 443 for HTTPS)
- Key Features:
 - Text-based protocol (easy to read & debug)
 - Stateless (each request is independent)
 - Supports multiple methods (GET, POST, etc.)



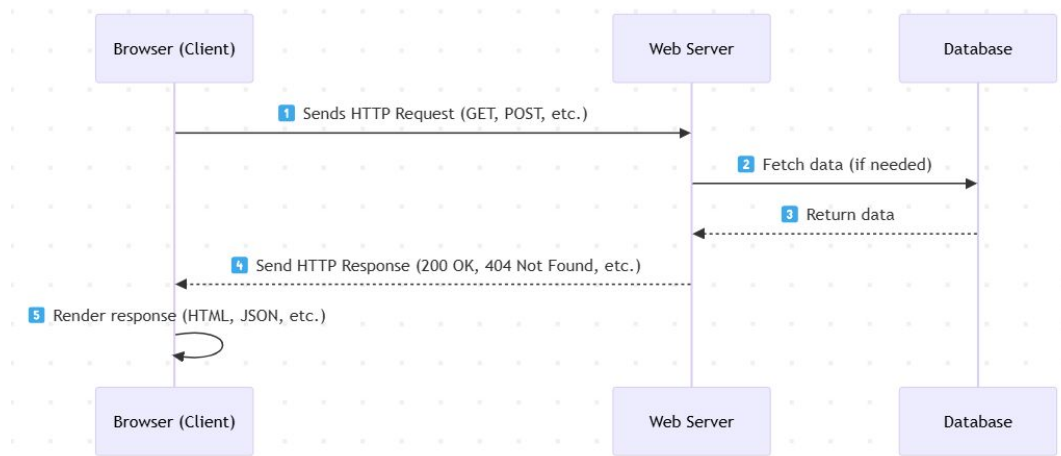
How HTTP Works

- Client-Server Model:
 - Client (browser or mobile app) makes an HTTP request
 - Server (web server, API, etc.) processes the request and sends back a response
- Components of an HTTP Request:
 - Method: Defines the action (GET, POST, etc.)
 - URL: The resource being requested
 - Headers: Metadata (e.g., user-agent, content type)
 - Body (optional): Data sent in POST/PUT requests
- Components of an HTTP Response:
 - Status Code: Indicates success or failure (e.g., 200 OK, 404 Not Found)
 - Headers: Metadata about the response
 - Body (optional): The actual content returned



The HTTP Request-Response Cycle

1. Step 1: The browser (client) sends a request
2. Step 2: The web server processes the request
3. Step 3: The server generates a response and sends it back
4. Step 4: The browser renders the response (for a web page)



Stateless Nature of HTTP

- What does “stateless” mean?
 - HTTP does not retain memory of previous requests
 - Each request is treated as an independent transaction
- Challenges of Statelessness:
 - Hard to maintain user sessions (e.g., login state)
 - Each request must carry all necessary information
- How do we handle state?
 - Cookies – Small pieces of data stored in the browser
 - Sessions – Server-side storage of user state
 - Tokens (JWT, OAuth) – Used for authentication & authorization

HTTP Methods

- Common HTTP Methods & Their Use Cases:
 - GET: Retrieve a resource (e.g., webpage, API data)
 - POST: Send data to create a new resource (e.g., form submission)
 - PUT: Update an existing resource
 - DELETE: Remove a resource
 - PATCH: Partially update a resource

HTTP Status Codes

- 1xx – Informational (Request received, continuing process)
- 2xx – Success (Request successfully processed)
 - 200 OK – Successful response
 - 201 Created – Resource successfully created
- 3xx – Redirection (Further action needed)
 - 301 Moved Permanently – Resource URL changed
 - 304 Not Modified – Use cached version
- 4xx – Client Errors (Mistakes in the request)
 - 400 Bad Request – Incorrect request format
 - 401 Unauthorized – Authentication required
 - 403 Forbidden – No permission
 - 404 Not Found – Resource doesn't exist
- 5xx – Server Errors (Issue with the server)
 - 500 Internal Server Error – Unexpected server failure
 - 503 Service Unavailable – Server overloaded

What About HTTPS?

- HTTPS is the secure version of HTTP, using SSL/TLS encryption.
- It ensures data confidentiality, integrity, and authentication.
- Used for secure websites, online banking, and e-commerce.
- Works on Port 443 instead of Port 80.

HTTP Interview Questions

- Basic Questions:
 - What is HTTP, and how does it work?
 - Why is HTTP considered a stateless protocol?
 - What are the key differences between HTTP and HTTPS?
- Intermediate Questions:
 - Explain the HTTP request-response cycle with an example.
 - What are HTTP methods? When would you use PUT vs. PATCH?
 - What are HTTP status codes? Give examples of 2xx, 3xx, 4xx, and 5xx status codes.
- Advanced Questions:
 - How do cookies, sessions, and tokens help maintain state in HTTP?
 - What is the difference between 301 (Moved Permanently) and 302 (Found) redirections?
 - How does caching work in HTTP, and which headers control it?
 - What security risks are associated with HTTP, and how can they be mitigated?

Summary & Key Takeaways


- HTTP is the foundation of web communication
- It follows a request-response model
- HTTP is stateless, requiring mechanisms like cookies & sessions
- HTTP methods define actions, and status codes indicate responses
- What's Next:
 - REST & RESTfulness - API Design Principles

REST & RESTfulness - API Design Principles

— Protocols —

What is REST?

- **Definition:** REST (Representational State Transfer) is an architectural style for designing networked applications.
- **Key Idea:** Uses standard HTTP methods and stateless communication.
- **Origin:** Coined by Roy Fielding in his 2000 dissertation.



Why REST Matters?

- **Simplicity & Scalability:** Based on standard HTTP.
- **Interoperability:** Works across different platforms.
- **Efficiency:** Uses caching, statelessness for performance.

REST Constraints (Core Principles)

- Client-Server Architecture
- Statelessness
- Cacheability
- Layered System
- Uniform Interface

Client-Server

Cache

Stateless

Uniform Contract

Layered System

Code-on-Demand

RESTful API Design Principles

- Resource-Based Approach
 - GET `/users/{id}` to retrieve a user
 - POST `/orders` to create a new order
- Proper HTTP Methods Usage
 - GET to **retrieve** data
 - POST to **create** new resources
 - PUT to **update** existing resources
 - PATCH for **partial updates**
 - DELETE to **remove** resources
- Stateless Interactions
- Consistency in URL Structure
 - ☒ Use **plural nouns** for collections: `/users`, `/orders`
 - ☒ Avoid including actions in URLs: `/users/{id}/activate` ❌ → `/users/{id}` with PATCH ☒
 - ☒ Implement **versioning** for backward compatibility: `/v1/users`

Resources & Endpoints

- Resources: Entities like Users, Orders, Products.
- Endpoint Examples:
 - GET /users/{id}
 - POST /orders
 - DELETE /products/{id}

GET `https://example.com/products` → Get all products

GET `https://example.com/products/{id}` → Get details of a specific product

POST `https://example.com/products` → Add a new product

PUT `https://example.com/products/{id}` → Update a product

PATCH `https://example.com/products/{id}` → Partially update a product

DELETE `https://example.com/products/{id}` → Remove a product

JSON vs. XML in REST APIs

- Why JSON?
 - Lightweight
 - Faster parsing
 - Readable
- When to Use XML?
 - Legacy systems
 - Data validation needs

{JSON}



<XML>

HTTP Methods in REST

- **GET**: Retrieve data
- **POST**: Create a resource
- **PUT**: Update a resource
- **PATCH**: Partial update
- **DELETE**: Remove a resource
- **OPTIONS, HEAD**: Additional functionalities

GET `https://example.com/orders` → Fetch all orders

GET `https://example.com/orders/{id}` → Fetch a specific order

POST `https://example.com/orders` → Create a new order

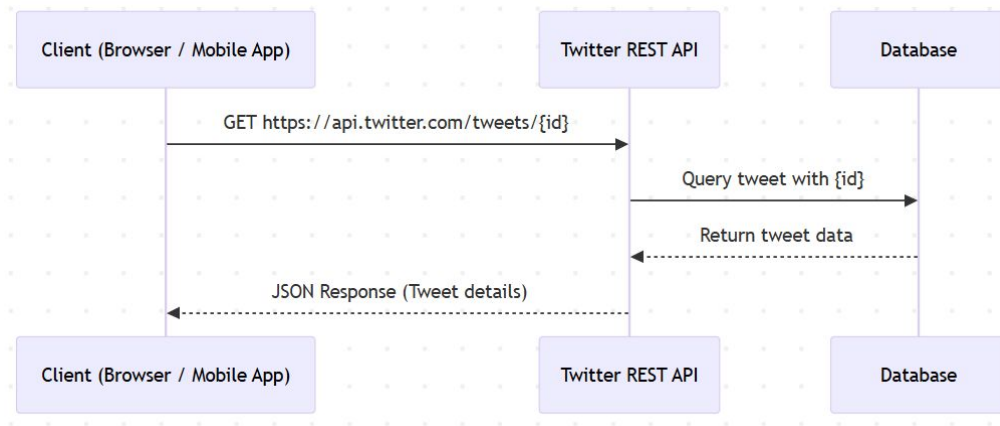
PUT `https://example.com/orders/{id}` → Update an order

PATCH `https://example.com/orders/{id}` → Partially update an order

DELETE `https://example.com/orders/{id}` → Cancel an order


Real-World REST API Examples

- Twitter API
 - Fetch tweets: GET /tweets/{id}
 - Post a tweet: POST /tweets
- GitHub API
 - Get repo details: GET /repos/{owner}/{repo}
 - Create an issue: POST /repos/{owner}/{repo}/issues



Best Practices & Common Pitfalls

- ✓ Use Proper Status Codes (200, 201, 400, 404, 500)
- ✓ Versioning APIs (/v1/resource)
- ✓ Implement Authentication (OAuth, JWT)
- ✓ Pagination (?page=2&limit=20)
- ✗ Avoid using Verbs in URLs (/createUser → POST /users)



Interview Questions - REST

- Basic Questions
 - What is REST, and how does it differ from SOAP?
 - What are the six constraints of REST architecture?
 - What is the difference between a REST API and a RESTful API?
 - What is a resource in REST, and how is it represented?
 - What are endpoints in a REST API?
- HTTP Methods & Status Codes
 - Explain the difference between GET, POST, PUT, PATCH, and DELETE.
 - When would you use PUT vs. PATCH?
 - What are the commonly used HTTP status codes in REST APIs?
 - What does the 200 OK, 201 Created, 204 No Content, 400 Bad Request, 401 Unauthorized, 403 Forbidden, 404 Not Found, and 500 Internal Server Error mean?
- RESTful API Design & Best Practices
 - What are the best practices for designing RESTful APIs?
 - How do you design a RESTful API for a blogging platform (e.g., users, posts, comments)?
 - What is HATEOAS (Hypermedia as the Engine of Application State) in REST?
 - How do you handle authentication and authorization in REST APIs?
 - What are API rate limiting and throttling, and why are they important?
 - How do you ensure backward compatibility when updating a REST API?
- Advanced & Real-World Questions
 - How does caching work in RESTful APIs?
 - What is the difference between JSON and XML in REST?
 - How do you implement pagination in REST APIs?
 - How does versioning work in REST APIs?
 - Explain the differences between REST, GraphQL, and gRPC.
 - How would you improve the performance of a REST API?
 - Can a REST API be stateful? Why or why not?
 - How does REST handle security vulnerabilities like SQL injection and CSRF?

Summary & Key Takeaways

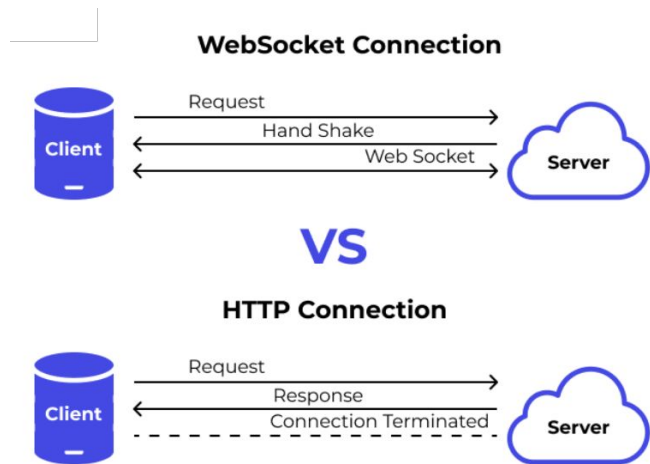
- REST is a stateless, scalable, and widely adopted API design style.
- Proper endpoint structure, HTTP methods, and response formats matter.
- Security, versioning, and best practices improve API quality.
- What's next:
 - Real-Time Communication Protocols

Real-Time Communication Protocols

— Protocols —


Introduction to Real-Time Communication

- **Definition:** Real-time communication refers to the continuous exchange of data with minimal latency.
- Why real-time is important (e.g., instant chat, live stock updates, multiplayer games).
- Challenges of traditional request-response HTTP model.
- Traditional HTTP vs. Real-Time Communication
- Alternatives to improve real-time data exchange:
 - Polling
 - WebSockets
 - Server-Sent Events (SSE)
 - Long Polling



WebSockets: Persistent Full-Duplex Communication

- Definition: WebSockets provide a persistent, full-duplex connection between the client and server over a single TCP connection.
- How they work:
 - WebSocket handshake using HTTP upgrade request.
 - **Step 1:** Client requests an upgrade to WebSockets
 - **Step 2:** Server accepts and keeps the connection open.
 - Connection remains open for continuous data exchange.
 - Either client or server can send messages at any time.
 - **Step 3:** Data is exchanged in real-time using frames.
 - **Step 4:** Either party can close the connection when done.




Advantages and use cases of WebSockets

- Advantages:
 - Persistent connection = lower latency.
 - Reduces overhead compared to HTTP polling.
 - Efficient for real-time applications.
- Use cases:
 - Live Chat Applications (WhatsApp, Slack, Discord)
 - Stock Market Price Updates (NASDAQ, Robinhood)
 - Multiplayer Online Games (Fortnite, Call of Duty)
 - Collaborative Tools (Google Docs, Figma)







Long Polling: Simulating Real-Time with HTTP

- **Definition:** A technique where the client sends a request to the server and waits until the server has new data to respond with.
- How it differs from regular polling:
- Instead of immediate responses, the server holds the request until new data is available.
- How Long Polling Works (Step-by-Step)
 - a. Client makes an HTTP request.
 - b. Server holds the request until data is available.
 - c. Server responds with new data.
 - d. Client immediately sends another request.



When to Use WebSockets vs. Long Polling?

- Use WebSockets when:
 -  High-frequency, bi-directional data exchange is needed.
 -  Low latency is critical (e.g., gaming, chat).
- Use Long Polling when:
 -  WebSockets are not supported or overkill.
 -  Periodic updates are sufficient (e.g., notifications).
- Real-world example:
 - Slack: WebSockets for chat.
 - Twitter: Long Polling for notifications.
 - Stock Exchanges: WebSockets for real-time data feeds.
 - IoT devices: Long Polling for intermittent updates.

Interview Questions

- Basic Questions

- What is real-time communication, and why is it important?
- How does WebSockets work, and how does it differ from traditional HTTP?
- Explain the WebSocket handshake process.
- What is long polling, and how does it work?

- Intermediate Questions

- What are the advantages of WebSockets over long polling?
- In what scenarios would you prefer long polling over WebSockets?
- How does WebSockets handle connection failures or network interruptions?
- Can you use WebSockets with load balancers? If yes, how?
- What are some challenges of scaling WebSockets in a distributed system?

Summary & Final Takeaways


- WebSockets = Persistent, full-duplex communication.
- Long Polling = Simulated real-time via HTTP requests.
- Choose based on latency needs and infrastructure.
- What's next:
 - Modern API Protocols - Beyond REST(gRPC, GraphQL)

Modern API Protocols - Beyond REST(gRPC, GraphQL)

— Protocols —

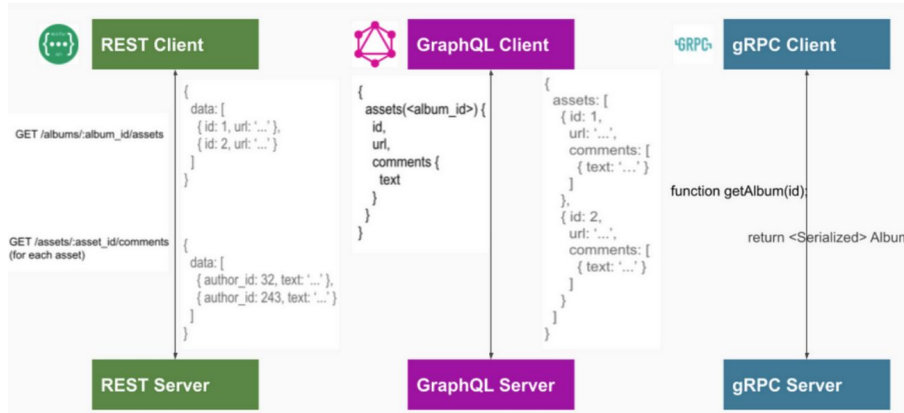
Why Do We Need More Than REST?

- Limitations of REST APIs:
 - Over-fetching & Under-fetching (Clients get too much or too little data).
 - High Latency (Multiple requests needed for complex data).
 - Not optimized for real-time communication (Polling required).
- Need for modern solutions like gRPC & GraphQL to handle efficiency, flexibility, and performance.



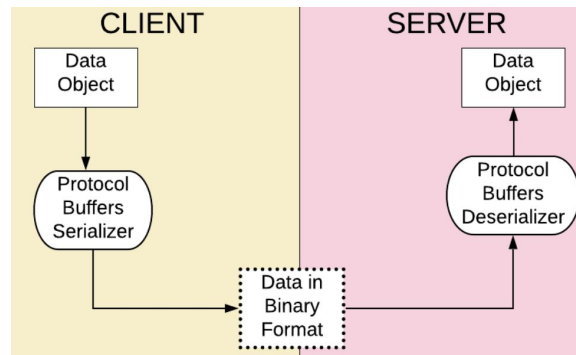
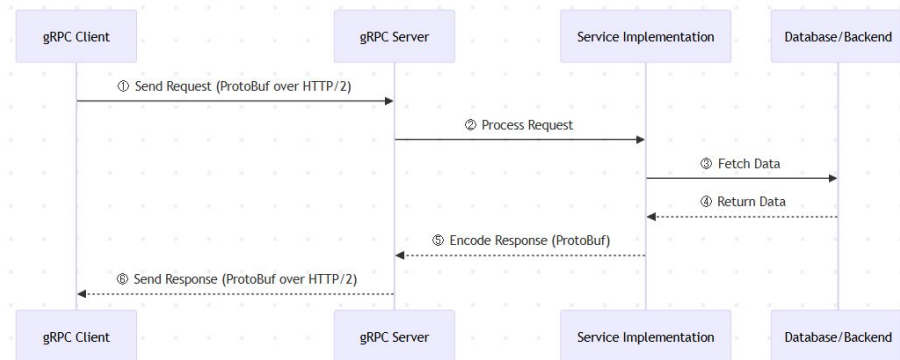
Introducing gRPC & GraphQL

- **gRPC:** A high-performance, binary protocol optimized for microservices & real-time communication.
- **GraphQL:** A flexible query language that allows clients to fetch only the data they need.
- Key Differences:
 - gRPC = Fast, binary, microservices-focused
 - GraphQL = Flexible, JSON-based, frontend-optimized



gRPC - How It Works?

- gRPC is built on HTTP/2, allowing:
 - Multiplexed requests (multiple calls over one connection).
 - Compression (smaller payload sizes).
 - Full-duplex streaming (real-time bidirectional communication).
- Uses Protocol Buffers (ProtoBuf) for fast serialization (smaller & faster than JSON).

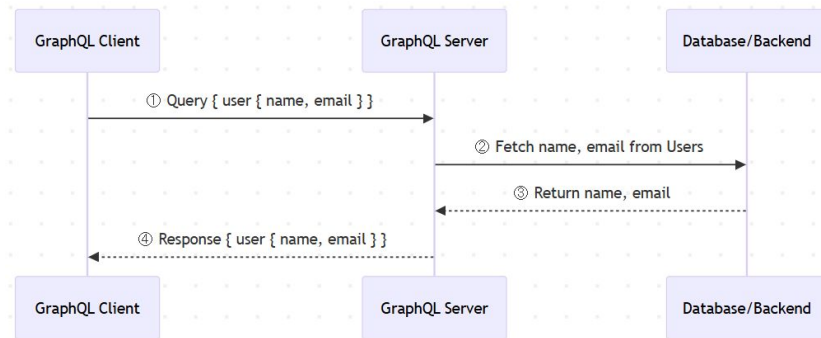


gRPC - Use Cases & When to Use

- Microservices – Fast inter-service communication.
- Real-time streaming – Full-duplex bidirectional data transfer.
- IoT & Low-bandwidth applications – Efficient binary communication.
- Multi-language ecosystems – Auto-generated client & server code.

GraphQL - How It Works?

- Instead of multiple REST endpoints, GraphQL has one endpoint where clients specify the data they need.
- GraphQL Schema defines types & relationships between data.
- Clients send queries → GraphQL server resolves fields dynamically.



```
query {  
  user(id: 123) {  
    name  
    email  
  }  
}
```

```
{  
  "data": {  
    "user": {  
      "name": "John Doe",  
      "email": "john.doe@example.com"  
    }  
  }  
}
```

GraphQL - Use Cases & When to Use

- Frontend Optimization – Clients fetch exactly what they need.
- Reducing API Requests – One query replaces multiple REST calls.
- Mobile & Web Apps – Handles slow networks & multiple data sources.
- Aggregating Data from Multiple Services – Simplifies fetching data from different databases or APIs.

Interview Questions on API Protocols

- How would you compare REST, gRPC, and GraphQL?
- When would you use gRPC over REST?
- What are the trade-offs of using GraphQL in a large-scale system?
- How does gRPC handle authentication & security?
- How do you scale GraphQL APIs efficiently?

Summary & Key Takeaways

- gRPC → Best for microservices, real-time streaming, high-performance APIs.
- GraphQL → Best for frontend-driven APIs, flexible data fetching, and reducing over-fetching.
- Choose based on use case – No one-size-fits-all!
- Be ready to justify your API choices in system design interviews!
- What's next:
 - Summary & Practical Applications of protocols

Summary - Protocols

- In this section, we explored key communication and API design protocols essential for system design:
 - TCP & UDP - Understanding the Basics
 - HTTP - The Backbone of the Web
 - REST & RESTfulness - API Design Principles
 - Real-Time Communication Protocols
 - Modern API Protocols - Beyond REST
- Next Up:
 - Architectural Patterns
 - We now shift focus to architectural patterns, which define how components interact in scalable systems.